


CSC 261/461

Eustrat Zhupa



Structured Query Language - INSERT

- ▶ **INSERT** is used to add a single tuple to a table
 - ▶ have to specify the table name and a list of values for the tuple.
 - ▶ values should be listed in the same order in which they were defined with **CREATE TABLE**.

```
INSERT INTO EMPLOYEE  
VALUES ('Richard','K','Marini','653298653',  
'1962-12-30','98 Oak Forest, Katy, TX', 'M',  
37000, '653298653', 4 );
```

Structured Query Language - INSERT

- ▶ the user can specify explicit attribute names that correspond to the values provided in the INSERT command.
- ▶ useful if a relation has many attributes but you assign only a few.
- ▶ the values must include all attributes with NOT NULL specification and no default value.
- ▶ Attributes with NULL or DEFAULT values are the ones that can be *left out*.

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Richard', 'Marini', 4, '653298653');
```

Structured Query Language - INSERT

You can also add many tuples at a time

```
CREATE TABLE WORKS_ON_INFO
```

```
( Emp_name VARCHAR(15),  
  Proj_name VARCHAR(15),  
  Hours_per_week DECIMAL(3,1) );
```

```
INSERT INTO WORKS_ON_INFO ( Emp_name, Proj_name, Hours_per_week )  
SELECT E.Lname, P.Pname, W.Hours  
FROM PROJECT P, WORKS_ON W, EMPLOYEE E  
WHERE P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```

Structured Query Language - DELETE

- ▶ **DELETE** removes tuples from a relation.
- ▶ includes a **WHERE** clause to select the tuples to be deleted.
- ▶ Tuples are explicitly deleted from only one table at a time.
 - ▶ may propagate to other relations if referential triggered actions are specified.
- ▶ a missing **WHERE** deletes all tuples in the relation
 - ▶ table remains in the database as an empty table.

Structured Query Language - DELETE

```
DELETE FROM EMPLOYEE  
WHERE Lname='Brown';
```

```
DELETE FROM EMPLOYEE  
WHERE Ssn='123456789';
```

```
DELETE FROM EMPLOYEE  
WHERE Dno=5;
```

```
DELETE FROM EMPLOYEE;
```

[md]

Structured Query Language - Example

```
UPDATE PROJECT  
SET Plocation = 'Bellaire', Dnum = 5  
WHERE Pnumber=10;
```

```
UPDATE EMPLOYEE  
SET Salary = Salary * 1.1  
WHERE Dno = 5;
```

Note: *Old* salary on the right, *new* salary on the left.

[mu]

Structured Query Language - NULLs

- ▶ SQL has various rules for dealing with NULL values.
- ▶ NULL values represent:
 1. Unknown value.
 2. Unavailable or withheld value.
 3. Not applicable attribute.
- ▶ SQL *does not distinguish* between these cases.

Structured Query Language

Table 5.1 Logical Connectives in Three-Valued Logic

| | | | | |
|-----|------------|---------|---------|---------|
| (a) | AND | TRUE | FALSE | UNKNOWN |
| | TRUE | TRUE | FALSE | UNKNOWN |
| | FALSE | FALSE | FALSE | FALSE |
| | UNKNOWN | UNKNOWN | FALSE | UNKNOWN |
| (b) | OR | TRUE | FALSE | UNKNOWN |
| | TRUE | TRUE | TRUE | TRUE |
| | FALSE | TRUE | FALSE | UNKNOWN |
| | UNKNOWN | TRUE | UNKNOWN | UNKNOWN |
| (c) | NOT | | | |
| | TRUE | FALSE | | |
| | FALSE | TRUE | | |
| | UNKNOWN | UNKNOWN | | |

Structured Query Language

- ▶ SQL allows queries that check whether an attribute value is NULL.
- ▶ to compare an attribute value to NULL, use **IS** or **IS NOT**.
- ▶ **Query** Retrieve names (first, last) of all employees who do *not* have supervisors.

```
SELECT Fname, Lname  
FROM EMPLOYEE  
WHERE Super_ssn IS NULL;
```

Nested Queries

Sometimes you need to compare against a (multi)set of values (tuples).

- ▶ SQL uses the comparison operator **IN** to compare a value v with a set (or multiset) of values V
 - ▶ evaluates to TRUE if v is in V .
- ▶ If a single attribute in a single tuple is returned, the query result will be a single value.
 - ▶ you can use $=$ instead of **IN**.

Q4: Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

[2]

Structured Query Language

Q4A: (**SELECT** **DISTINCT** Pnumber
FROM PROJECT, DEPARTMENT, EMPLOYEE
WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn
AND Lname='Smith')

```
UNION
( SELECT DISTINCT Pnumber
  FROM PROJECT, WORKS_ON, EMPLOYEE
 WHERE Pnumber=Pno AND Essn=Ssn
    AND Lname='Smith' );
```

```

Q4A:  SELECT      DISTINCT Pnumber
      FROM        PROJECT
      WHERE       Pnumber IN
                ( SELECT      Pnumber
                  FROM        PROJECT, DEPARTMENT, EMPLOYEE
                  WHERE       Dnum=Dnumber AND
                             Mgr_ssn=Ssn AND Lname='Smith' )

      OR

      Pnumber IN
      ( SELECT      Pno
        FROM        WORKS_ON, EMPLOYEE
        WHERE       Essn=Ssn AND Lname='Smith' );

```

Nested Queries

- ▶ SQL allows the use of tuples of values in comparisons by placing them within parentheses.

```
SELECT DISTINCT Essn
FROM WORKS_ON
WHERE (Pno, Hours) IN ( SELECT Pno, Hours
                        FROM WORKS_ON
                        WHERE Essn='123456789' );
```

Other Operators

Other comparison operators can be used to compare a single value v to a set V .

- ▶ The = **ANY** (**SOME**) operator returns TRUE if the value v is equal to some value in the set V and is hence equivalent to **IN**.
- ▶ Other operators that can be combined with **ANY** include $>$, $>=$, $<$, $<=$, and $<>$.
- ▶ **ALL** can also be combined with each of these operators.

```
SELECT      Lname, Fname
FROM        EMPLOYEE
WHERE       Salary > ALL ( SELECT      Salary
                           FROM        EMPLOYEE
                           WHERE       Dno=5 );
```

[a]

Correlated Nested Queries

Nested query condition references attributes of the outer query.

- ▶ **EXISTS** and **NOT EXISTS** are used in conjunction with a *correlated* nested query Q .
- ▶ **EXISTS(Q)** returns TRUE if there is at least one tuple in the result of Q , FALSE otherwise.
- ▶ **NOT EXISTS(Q)** returns TRUE if there are no tuples in the result of Q , FALSE otherwise.

Query: Retrieve names (first, last) of employees with no dependents.

NOT EXISTS

- ▶ **NOT EXISTS(Q)** returns TRUE if there are no tuples in the result Q, FALSE otherwise.

Query: Retrieve names (first, last) of employees with no dependents.

Query 6. Retrieve the names of employees who have no dependents.

```
Q6:      SELECT      Fname, Lname
          FROM        EMPLOYEE
          WHERE        NOT EXISTS ( SELECT      *
                                   FROM        DEPENDENT
                                   WHERE        Ssn=Essn );
```

Aggregate functions

Aggregate functions are used to summarize information from multiple tuples into a single-tuple summary.

- ▶ The **COUNT** function returns the number of tuples or values as specified in a query
- ▶ **SUM**, **MAX**, **MIN**, and **AVG** can be applied to a set (multi) of numeric values and return, respectively, the sum, maximum value, minimum value, and average (mean) of those values.

Aggregate Functions

Query: Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.

```
SELECT SUM(Salary), MAX(Salary), MIN(Salary), AVG(Salary)
FROM EMPLOYEE;
```

Query: Find the sum of the salaries of all employees of the Research department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
SELECT SUM(Salary), MAX(Salary), MIN(Salary), AVG(Salary)
FROM EMPLOYEE, DEPARTMENT
WHERE Dname = 'Research';
```

Grouping

- ▶ In many cases we want to apply the *aggregate* functions to subgroups of tuples in a relation, with subgrouping based on some attribute values.
 - ▶ find the *average* salary of employees in each department.
- ▶ we need to partition the relation into *disjoint* subsets of tuples.
- ▶ each group consist of the tuples with the same value of some attribute(s), called *grouping* attribute(s).

Query 24. For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
Q24:  SELECT    Dno, COUNT (*), AVG (Salary)
      FROM      EMPLOYEE
      GROUP BY  Dno;
```

GROUP BY

(a)

| Fname | Minit | Lname | Ssn | ... | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | ... | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | | 40000 | 888665555 | 5 |
| Ramesh | K | Narayan | 666884444 | | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | | 25000 | 333445555 | 5 |
| Alicia | J | Zelaya | 999887777 | | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | | 43000 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | | 25000 | 987654321 | 4 |
| James | E | Bong | 888665555 | | 55000 | NULL | 1 |

| Dno | Count (*) | Avg (Salary) |
|-----|-----------|--------------|
| 5 | 4 | 33250 |
| 4 | 3 | 31000 |
| 1 | 1 | 55000 |

Result of Q24

Grouping EMPLOYEE tuples by the value of Dno

HAVING

Query 26. For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

```
Q26:  SELECT      Pnumber, Pname, COUNT (*)
      FROM        PROJECT, WORKS_ON
      WHERE       Pnumber=Pno
      GROUP BY    Pnumber, Pname
      HAVING      COUNT (*) > 2;
```

| Pname | Pnumber | ... | Essn | Pno | Hours |
|-----------------|---------|-----|-----------|-----|-------|
| ProductX | 1 | | 123456789 | 1 | 32.5 |
| ProductX | 1 | | 453453453 | 1 | 20.0 |
| ProductY | 2 | | 123456789 | 2 | 7.5 |
| ProductY | 2 | | 453453453 | 2 | 20.0 |
| ProductY | 2 | | 333445555 | 2 | 10.0 |
| ProductZ | 3 | | 666884444 | 3 | 40.0 |
| ProductZ | 3 | | 333445555 | 3 | 10.0 |
| Computerization | 10 | | 333445555 | 10 | 10.0 |
| Computerization | 10 | | 999887777 | 10 | 10.0 |
| Computerization | 10 | | 987987987 | 10 | 35.0 |
| Reorganization | 20 | | 333445555 | 20 | 10.0 |
| Reorganization | 20 | | 987654321 | 20 | 15.0 |
| Reorganization | 20 | | 888665555 | 20 | NULL |
| Newbenefits | 30 | | 987987987 | 30 | 5.0 |
| Newbenefits | 30 | | 987654321 | 30 | 20.0 |
| Newbenefits | 30 | | 999887777 | 30 | 30.0 |

- These groups are not selected by the HAVING condition of Q26.

After applying the WHERE clause but before applying HAVING

HAVING

Query 26. For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

```
Q26:  SELECT  Pnumber, Pname, COUNT (*)
        FROM    PROJECT, WORKS_ON
        WHERE   Pnumber=Pno
        GROUP BY Pnumber, Pname
        HAVING  COUNT (*) > 2;
```

| Pname | Pnumber | ... | Essn | Pno | Hours |
|-----------------|---------|-----|-----------|-----|-------|
| ProductY | 2 | | 123456789 | 2 | 7.5 |
| ProductY | 2 | | 453453453 | 2 | 20.0 |
| ProductY | 2 | | 333445555 | 2 | 10.0 |
| Computerization | 10 | | 333445555 | 10 | 10.0 |
| Computerization | 10 | ... | 999887777 | 10 | 10.0 |
| Computerization | 10 | | 987987987 | 10 | 35.0 |
| Reorganization | 20 | | 333445555 | 20 | 10.0 |
| Reorganization | 20 | | 987654321 | 20 | 15.0 |
| Reorganization | 20 | | 888665555 | 20 | NULL |
| Newbenefits | 30 | | 987987987 | 30 | 5.0 |
| Newbenefits | 30 | | 987654321 | 30 | 20.0 |
| Newbenefits | 30 | | 999887777 | 30 | 30.0 |

| Pname | Count (*) |
|-----------------|-----------|
| ProductY | 3 |
| Computerization | 3 |
| Reorganization | 3 |
| Newbenefits | 3 |

Result of Q26
(Pnumber not shown)

After applying the HAVING clause condition

Questions?



Thank you!